

# Advertisement Click - Logistic Regression

February 3, 2022

## Advertisement Click - Logistic Regression

by Joseph Ramon

Will be working with an advertising data set, indicating whether or not a particular internet user clicked on an Advertisement.

Will try to create a model that will predict whether or not they will click on an ad based off the features of that user.

This data set contains the following features:

- Daily Time Spent on Site': consumer time on site in minutes
- Age': customer age in years
- Area Income': Avg. Income of geographical area of consumer
- Daily Internet Usage': Avg. minutes a day consumer is on the internet
- Ad Topic Line': Headline of the advertisement
- City': City of consumer
- Male': Whether or not consumer was male
- Country': Country of consumer
- Timestamp': Time at which consumer clicked on Ad or closed window
- Clicked on Ad': 0 or 1 indicated clicking on Ad

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

### 0.1 Get the Data

Read in the advertising.csv file and set it to a data frame called ad\_data.

```
[2]: ad_data = pd.read_csv('advertising.csv')
```

```
[3]: ad_data.head()
```

```
[3]:   Daily Time Spent on Site  Age  Area Income  Daily Internet Usage  \
0                68.95    35    61833.90                256.09
1                80.23    31    68441.85                193.77
2                69.47    26    59785.94                236.50
```

3	74.15	29	54806.18	245.89
4	68.37	35	73889.99	225.58

	Ad Topic Line	City	Male	Country \
0	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia
1	Monitored national standardization	West Jodi	1	Nauru
2	Organic bottom-line service-desk	Davidton	0	San Marino
3	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy
4	Robust logistical utilization	South Manuel	0	Iceland

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0
2	2016-03-13 20:35:42	0
3	2016-01-10 02:31:19	0
4	2016-06-03 03:36:18	0

```
[4]: ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Daily Time Spent on Site              1000 non-null   float64
1   Age                                    1000 non-null   int64
2   Area Income                           1000 non-null   float64
3   Daily Internet Usage                   1000 non-null   float64
4   Ad Topic Line                         1000 non-null   object
5   City                                   1000 non-null   object
6   Male                                   1000 non-null   int64
7   Country                               1000 non-null   object
8   Timestamp                             1000 non-null   object
9   Clicked on Ad                         1000 non-null   int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

```
[5]: ad_data.describe()
```

```
[5]:
```

	Daily Time Spent on Site	Age	Area Income \
count	1000.000000	1000.000000	1000.000000
mean	65.000200	36.009000	55000.000080
std	15.853615	8.785562	13414.634022
min	32.600000	19.000000	13996.500000
25%	51.360000	29.000000	47031.802500
50%	68.215000	35.000000	57012.300000
75%	78.547500	42.000000	65470.635000

```
max          91.430000    61.000000  79484.800000
```

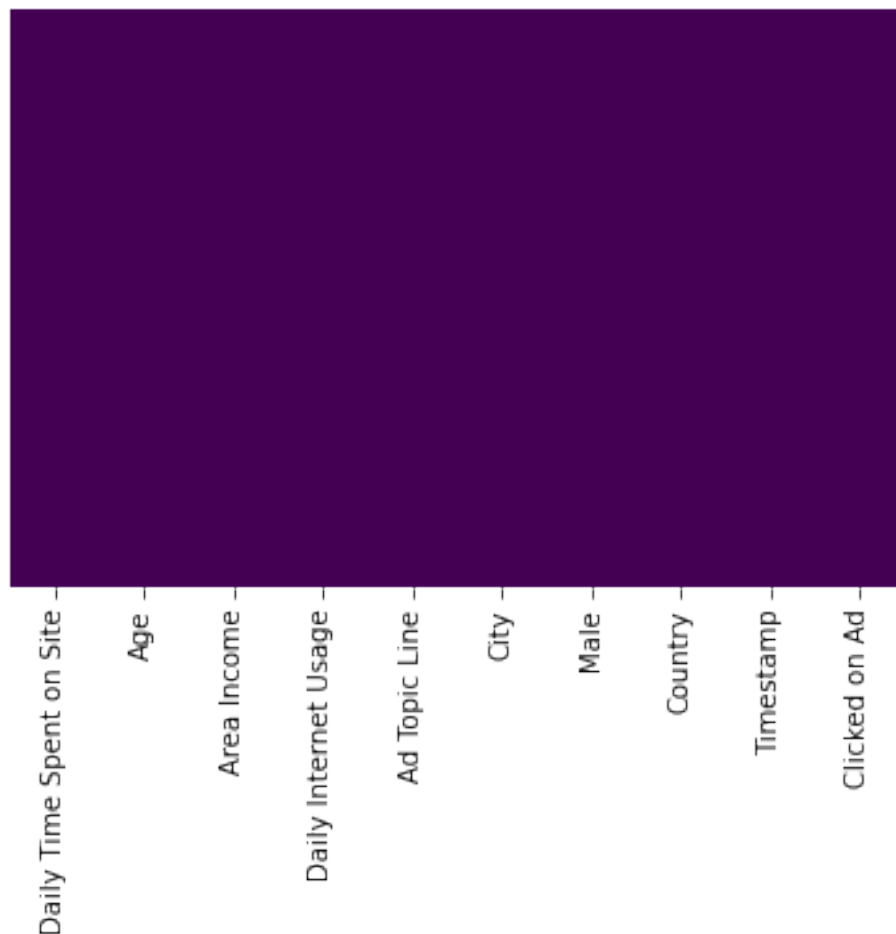
	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.000000
mean	180.000100	0.481000	0.500000
std	43.902339	0.499889	0.500250
min	104.780000	0.000000	0.000000
25%	138.830000	0.000000	0.000000
50%	183.130000	0.000000	0.500000
75%	218.792500	1.000000	1.000000
max	269.960000	1.000000	1.000000

## 0.2 Exploratory Data Analysis

### 0.3 Missing Data

```
[6]: sns.heatmap(ad_data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
     # the heatmap below does not show any nulls
```

```
[6]: <AxesSubplot:>
```



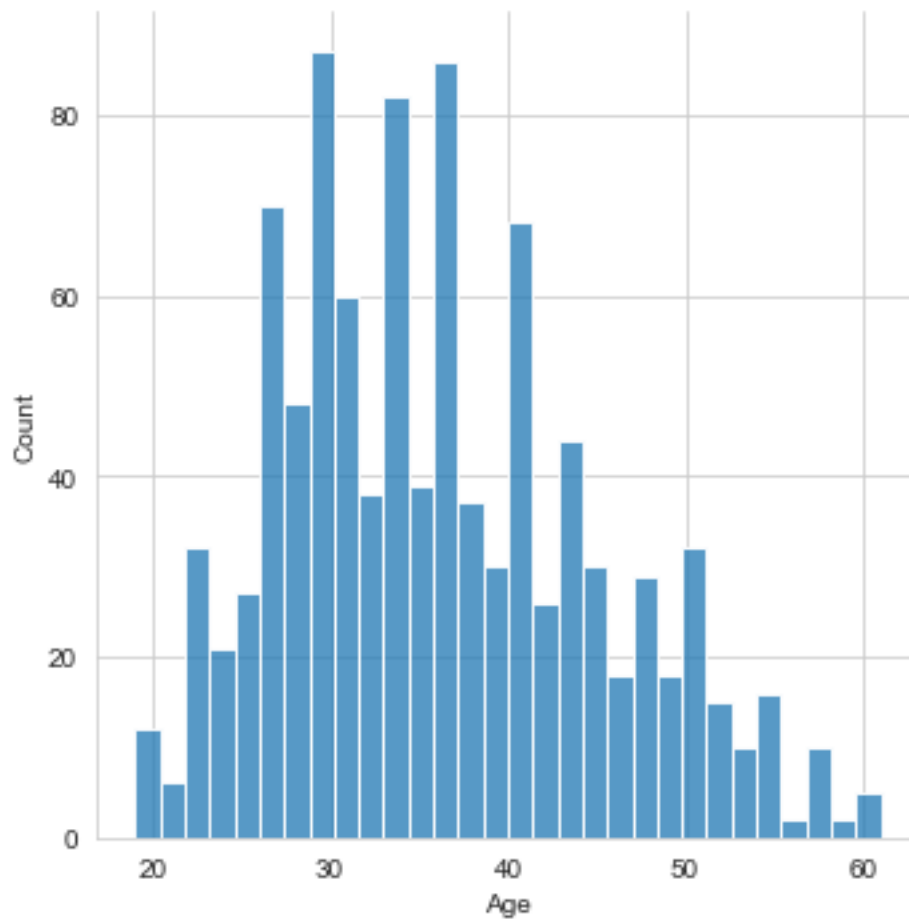
## 0.4 Some Plots

```
[7]: sns.set_style('whitegrid')
```

### Age Histogram - using seaborn

```
[41]: sns.displot(ad_data['Age'],bins=30)
```

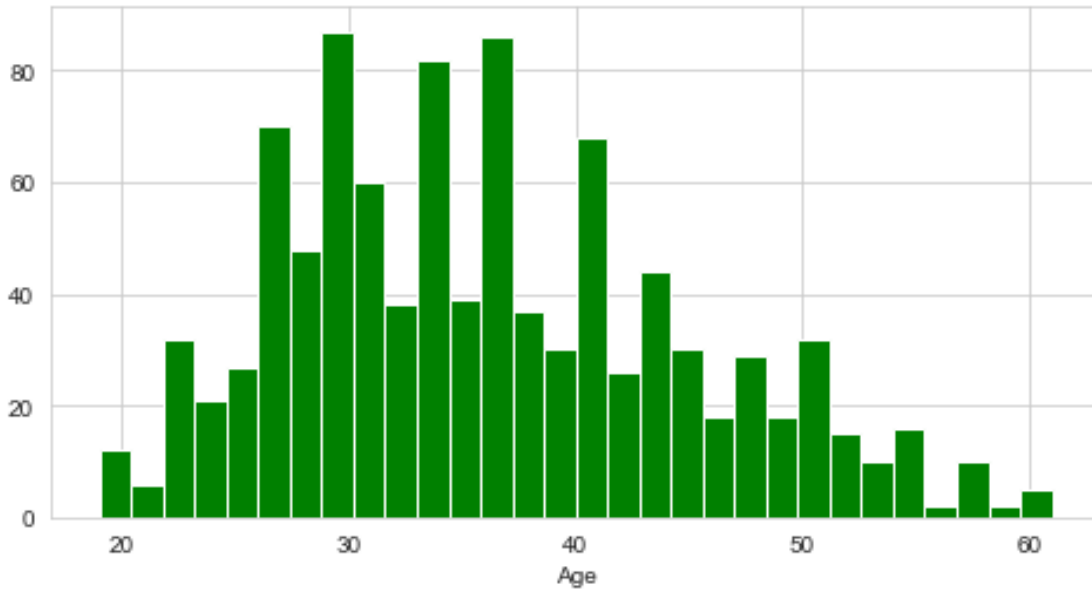
```
[41]: <seaborn.axisgrid.FacetGrid at 0x20e40861cd0>
```



### Age Histogram - using built-in

```
[37]: ad_data['Age'].hist(color='green',bins=30,figsize=(8,4))  
plt.xlabel('Age')
```

```
[37]: Text(0.5, 0, 'Age')
```



Age Histogram - interactive, using cufflinks


```
[86]: import cufflinks as cf
      cf.go_offline()
```

```
[87]: # more interactive version using cufflinks
      ad_data['Age'].iplot(kind='hist',bins=30,color='green',\
                           xTitle="Age", yTitle="Count", title="Age Count")
```

jointplot showing Area Income versus Age

```
[23]: sns.jointplot(x='Age',y='Area Income',data=ad_data,kind='scatter')
```

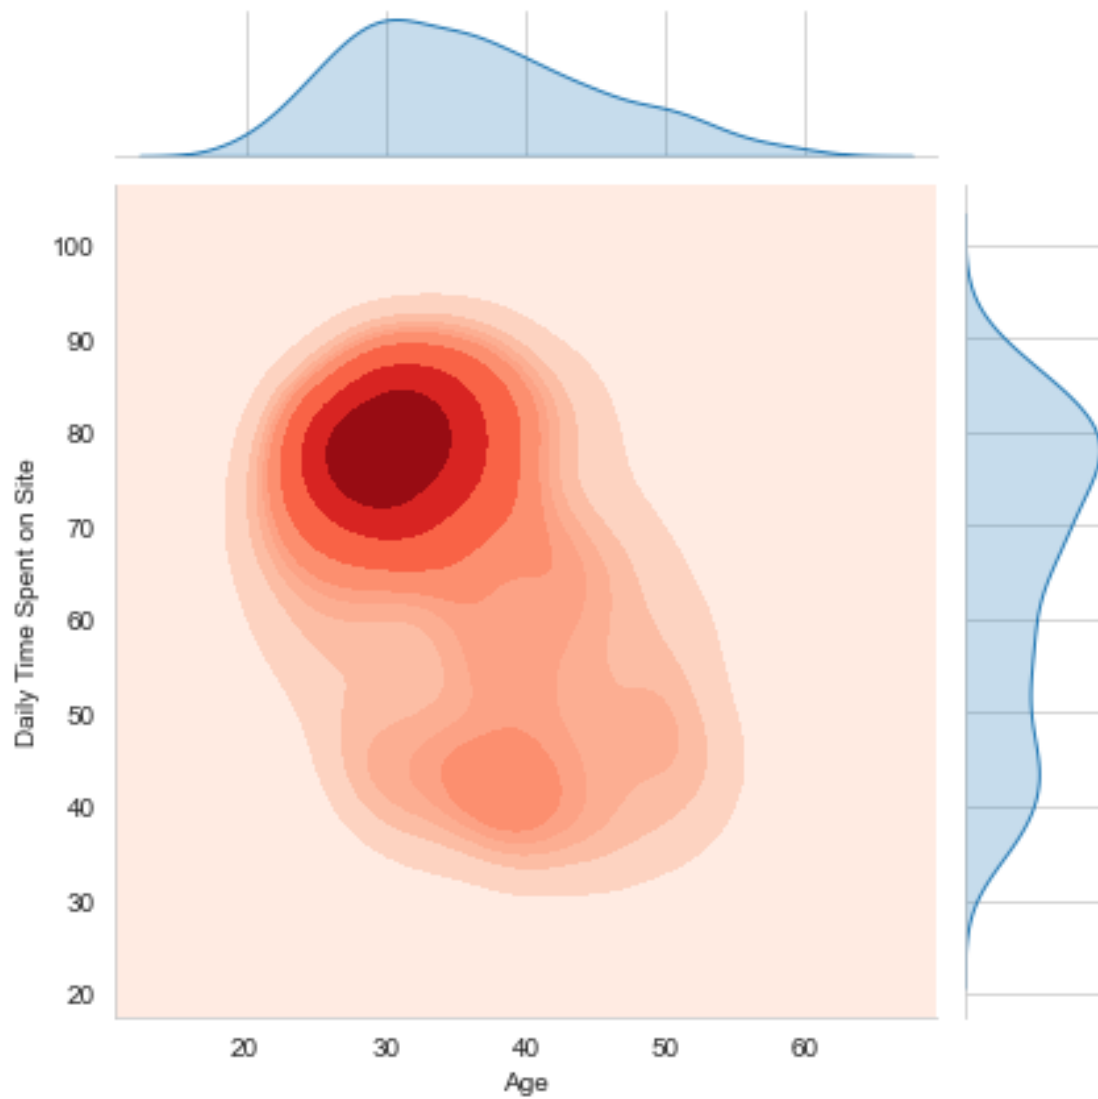
```
[23]: <seaborn.axisgrid.JointGrid at 0x20e3eff7d90>
```



jointplot showing the kde distributions of Daily Time spent on site vs. Age

```
[50]: sns.jointplot(x='Age',y='Daily Time Spent on Site',data=ad_data,\n                fill=True, thresh=0,cmap='Reds',kind='kde')
```

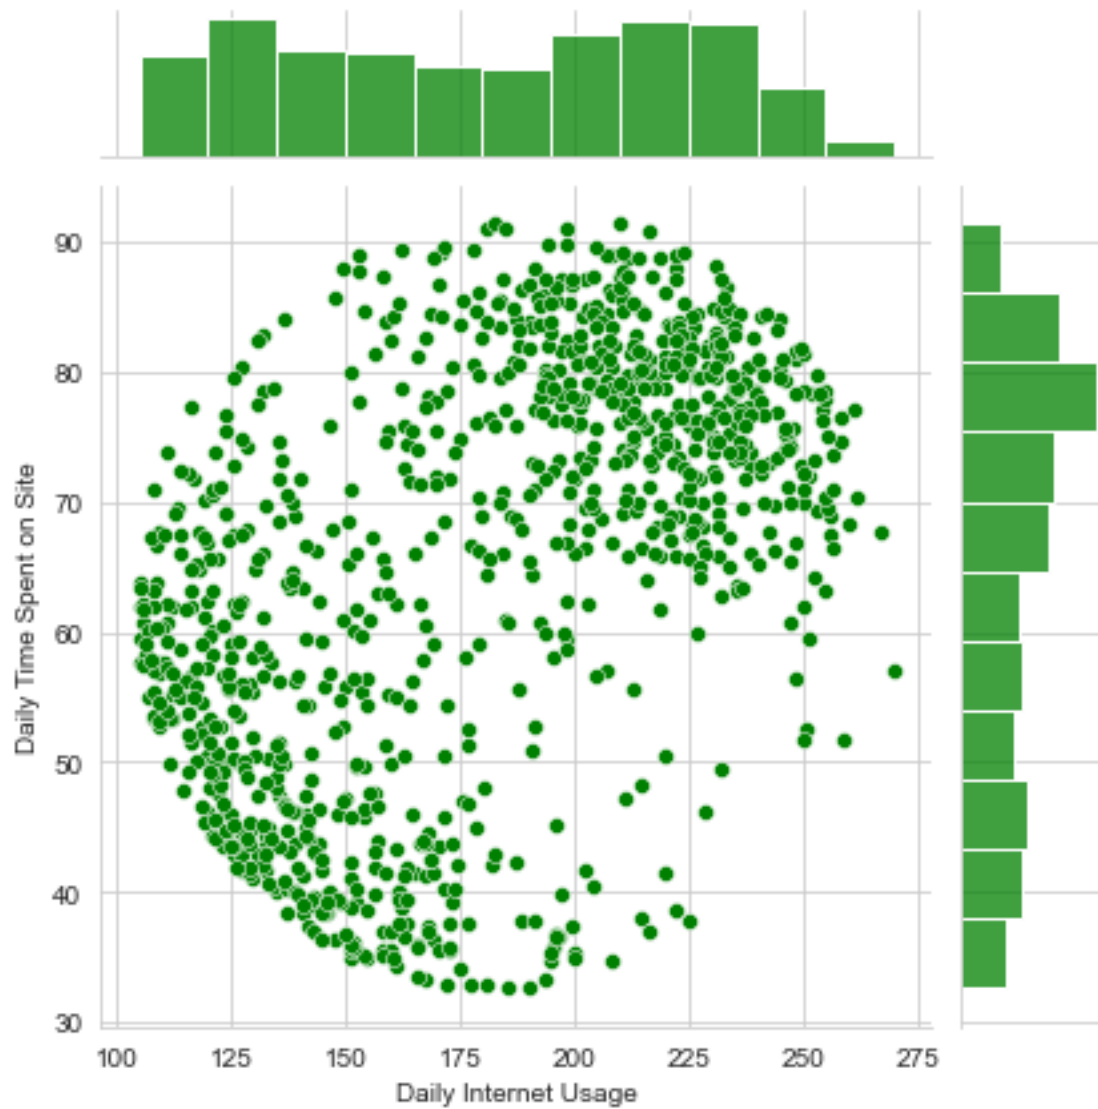
```
[50]: <seaborn.axisgrid.JointGrid at 0x20e41daa9a0>
```



jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage'

```
[55]: sns.jointplot(x='Daily Internet Usage',y='Daily Time Spent on Site',\n                  data=ad_data,color="Green")
```

```
[55]: <seaborn.axisgrid.JointGrid at 0x20e43344d30>
```

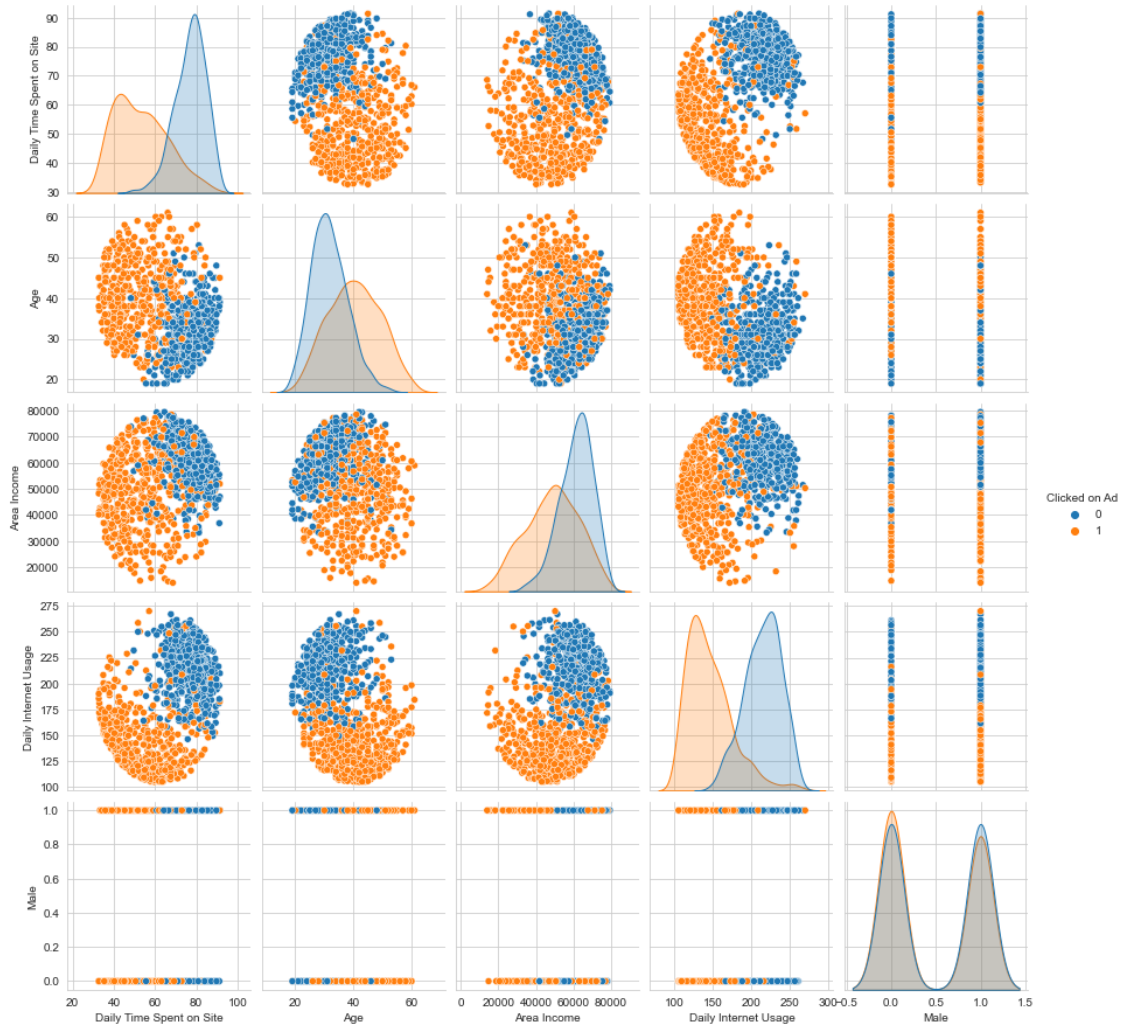


pairplot with the hue defined by the 'Clicked on Ad' column feature.

```
[58]: sns.pairplot(ad_data,hue='Clicked on Ad')
```

```
[58]: <seaborn.axisgrid.PairGrid at 0x20e447fe640>
```





## 1 Build a Logistic Regression model

Now it's time to do a train test split, and train our model

### 1.1 Train Test Split

```
[64]: ad_data.head(2)
```

```
[64]:   Daily Time Spent on Site  Age  Area Income  Daily Internet Usage  \
0                68.95    35    61833.90                256.09
1                80.23    31    68441.85                193.77
```

```
      Ad Topic Line      City  Male  Country  \
0  Cloned 5thgeneration orchestration  Wrightburgh    0  Tunisia
1  Monitored national standardization    West Jodi    1   Nauru
```

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0

```
[65]: ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Daily Time Spent on Site              1000 non-null   float64
1   Age                                    1000 non-null   int64
2   Area Income                           1000 non-null   float64
3   Daily Internet Usage                  1000 non-null   float64
4   Ad Topic Line                         1000 non-null   object
5   City                                   1000 non-null   object
6   Male                                   1000 non-null   int64
7   Country                               1000 non-null   object
8   Timestamp                             1000 non-null   object
9   Clicked on Ad                         1000 non-null   int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

```
[68]: ad_data.columns
```

```
[68]: Index(['Daily Time Spent on Site', 'Age', 'Area Income',
         'Daily Internet Usage', 'Ad Topic Line', 'City', 'Male', 'Country',
         'Timestamp', 'Clicked on Ad'],
         dtype='object')
```

```
[69]: from sklearn.model_selection import train_test_split
```

```
[76]: y = ad_data['Clicked on Ad']
      X = ad_data[['Daily Time Spent on Site', 'Age', 'Area Income',\
                  'Daily Internet Usage', 'Male']]
```

```
[77]: X_train, X_test, y_train, y_test = train_test_split(X, y,\
                                                         test_size=0.30, random_state=101)
```

**Train and fit a logistic regression model on the training set.**

```
[78]: from sklearn.linear_model import LogisticRegression
```

```
[79]: logmodel = LogisticRegression()
```

```
[80]: logmodel.fit(X_train,y_train)
```

```
[80]: LogisticRegression()
```

## 1.2 Predictions and Evaluations

predict values for the testing data

```
[94]: predictions = logmodel.predict(X_test)
```

Create a classification report for the model.

```
[83]: from sklearn.metrics import classification_report
```

```
[85]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.91	0.95	0.93	157
1	0.94	0.90	0.92	143
accuracy			0.93	300
macro avg	0.93	0.93	0.93	300
weighted avg	0.93	0.93	0.93	300